

MANVITH BELAME YOGESH

Phoenix, AZ (Open to relocation) | +1 617-318-8506 | manvith18@gmail.com | [LinkedIn](#) | [Portfolio](#) | [GitHub](#)

SUMMARY

Data engineer with 5+ years architecting multi-modal pipelines, data-quality frameworks, and lineage-tracked infrastructure in regulated environments. Proficient in Python, SQL, Airflow, and AWS, with a record of delivering research-ready datasets and validation systems that meet compliance standards.

PROFESSIONAL EXPERIENCE

Data Engineer, Walnutech AI, Chandler, AZ

Apr 2026 - Present

- Built the company's scholarship and pre-college discovery pipeline from scratch in Python and PostgreSQL (Supabase), automating extraction, enrichment, and ingestion to power a live catalog of **15,000+** scholarships and programs.
- Cut a critical field-extraction failure rate from **76%** to under **1%** by replacing hallucination-prone LLM extraction with deterministic logic validated against human-labeled records, ensuring students see accurate, working scholarship links.
- Re-engineered the monthly ingestion job to be idempotent and resumable with per-batch commits, trimming the crawl stage from **2.3 hours** to **10 minutes** and keeping the catalog refreshed reliably with no data loss on failures.
- Authored AI-native dev tooling as Claude Code agent skills, a self-revision review loop and a **0 to 100** production-readiness audit, wired into the PR workflow to catch data-integrity and migration risks before release.

Data Engineer, Community Dreams Foundation, Remote, AZ

Jun 2025 - Jun 2026

- Architected end-to-end ETL pipelines and a **9**-table Postgres data model with transactional batch-numbering, foreign keys, and indexing for a food-safety SaaS, enabling full data lineage across FDA-compliant worksheet outputs.
- Developed **9** JSON knowledge libraries as a RAG knowledge base, grounding Claude's hazard extraction across **10+** food products, sharply reducing worksheet hallucination and producing consistent, traceable outputs.
- Automated **50+** Python data-quality validation rules against JSON Schema definitions to catch structural regressions before deploying, hardening release reliability and keeping the FDA worksheet system audit-ready with zero production incidents.

Data Engineer Co-op, Plymouth Rock Assurance, Boston, MA

July 2024 - Jan 2025

- Migrated legacy SAS pipelines to Python using Polars and multithreading, converting SAS datasets to partitioned Parquet files on S3 and reducing pipeline runtime by **77%** with improved downstream query performance.
- Processed more than **300 million** rows of home insurance vendor data on AWS Redshift and S3, using Python and Polars, saving **80%** on cloud compute costs and built data validation frameworks to ensure data quality.
- Orchestrated end-to-end ETL workflows using AWS ECS Fargate and Step Functions, automating multi-source vendor data ingestion and improving operational efficiency by **60%** through cloud-native workflow orchestration.
- Identified and resolved critical data transformation errors corrupting quarterly records by implementing schema validation safeguards and reprocessing affected datasets to restore production data integrity.

Accenture Solutions, Bangalore, IN

Jan 2020 - Jan 2023

Senior Data Analyst

Jun 2022 - Jan 2023

- Conducted complex SQL-based ad-hoc analysis on customer behavior across six Disney applications, consolidating over **5 million** records to provide actionable insights to business stakeholders.
- Developed event-driven data pipelines with AWS EventBridge and Lambda to process real-time booking and system logs from AWS S3, enabling anomaly detection and quicker response to guest-impacting issues.
- Led a cross-functional team of **8** to deliver a high-stakes booking engine on schedule, enabling Disney to handle record visitor volumes flawlessly.
- Implemented automated park-pass booking validation using Python and SQL triggers, reducing manual processes by **80%**.

Data Analyst

Jan 2020 - May 2022

- Analyzed data from nearly half a million booking records using SQL to triage high-priority requests, accelerating response times and elevating the guest reservation experience.
- Built real-time dashboards to monitor peak load performance and ensure zero downtime during crucial ticket booking windows.
- Integrated Redshift data warehouse as a direct data source for Amazon QuickSight dashboards, boosting metric refresh speeds by **40%** to support faster executive decisions.
- Deployed CI/CD workflows on AWS, ensuring critical updates rolled out within **30 mins** after park-close, minimizing guest impact.

TECHNICAL SKILLS

Languages and databases: Python, Java, MySQL, MS SQL, MongoDB (NoSQL), PostgreSQL, Google BigQuery

Tools and Libraries: Tableau, Power BI, Jira, Git, MS Office Suite, Airflow, Splunk, Pandas, Polars, Matplotlib, PySpark

Cloud Computing: AWS (S3, Glue, Athena, SageMaker, RDS, Redshift, EMR), GCP, Microsoft Azure, Snowflake

Certifications: AWS Cloud Practitioner, Google Data Analytics Professional, Data Engineering Bootcamp (DataExpert.io)

EDUCATION

Northeastern University, Boston, MA

Jan 2023 - May 2025

Master of Science in Data Analytics Engineering

3.9 / 4.0 GPA

JSS Science and Technology University, Mysore, India

Aug 2016 - Jun 2020

Bachelor of Engineering